

Dirk Habich

Scalable Data Management (SDM)

Traditional Transaction Processing

Transactions

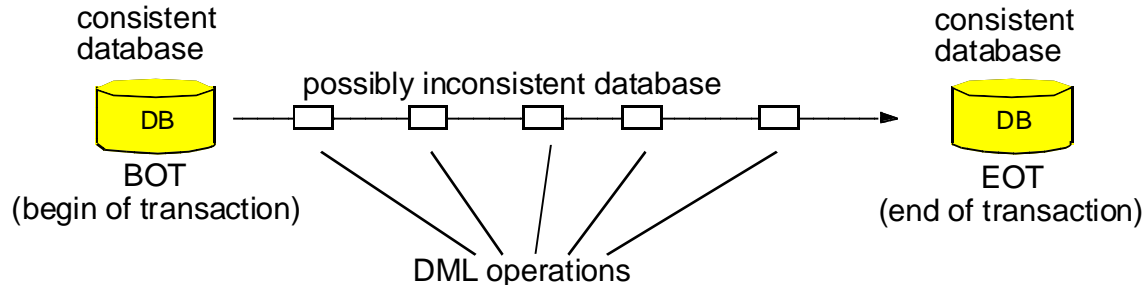
- A **transaction** is the execution of a **sequence** of one or more **operations** (e.g., SQL queries) on a shared database to perform some higher-level function
- It is the basic unit of change in a DBMS: Partial transactions are not allowed!

Example

- Move \$100 from Andy's bank account to his bookie's account
- Transaction
 - Check whether Andy has \$100
 - Deduct \$100 from his account
 - Add \$100 to his bookie's account

Principle of a transaction

- sequence of successive DB operations that transform a database from a consistent state into another consistent state
surrounded by: BOT EOT (Commit / Abort)



- properties
 - ACID: Atomicity, Consistency, Isolation, Durability
 - A transaction will always come to an end
 - Normal (commit): changes are permanently stored within the DB
 - Abnormal (abort/rollback): already composed changes are undone
- note: EOT state is not necessarily different from BOT state

Execute each transaction one-by-one (i.e., serial order) as they arrive at the DBMS

- One and only one transaction can be running at the same time in the DBMS

Before a transaction starts, copy the entire database to a new file and make all changes to that file

- If the transaction completes successfully, overwrite the original file with the new one
- If the transaction fails, just remove the dirty copy

→ Better approach: allow concurrent execution of independent transactions

- Question: Why do we want that?
 - Utilization/throughput (“hide” waiting for I/Os)
 - Increased response times to users
- But we also would like
 - Correctness
 - Fairness

Problem Statement

Arbitrary interleaving can lead to

- Temporary inconsistency (ok, unavoidable)
- Permanent inconsistency (bad!)

→ **Need formal correctness criteria**

Definitions

- A transaction may carry out many operations on the data retrieved from the database
- However, the DBMS is only concerned about what data is read/written from/to the database
 - Changes to the “outside world” are beyond the scope of the DBMS

Database

- A fixed set of named data objects (A, B, C, ...)

Transaction

- A sequence of read and write operations (R(A), W(B), ...)
 - DBMS's abstract view of a user program

Transactions in SQL

- A new transaction starts with the **begin** command
- The transaction stops with either **commit** or **abort**:
 - If **commit**, all changes are saved
 - If **abort**, all changes are undone so that it's like as if the transaction never executed at all
- Transaction can abort itself or the DBMS can abort it

Correctness Criteria: ACID

Atomicity

- All actions in the transaction happen, or none happen → *"all or nothing"*

Consistency

- If each transaction is consistent and the DB starts consistent, then it ends up consistent → *"it looks correct to me"*

Isolation

- Execution of one transaction is isolated from that of other transactions → *"as if alone"*

Durability

- If a transaction commits, its effects persist → *"survive failures"*



Recovery



Concurrency
Control

Atomicity of Transactions

Two possible outcomes of executing a transaction:

- Transaction might commit after completing all its actions
- or it could abort (or be aborted by the DBMS) after executing some actions

DBMS guarantees that transaction are atomic

- From user's point of view: transaction always either executes all its actions, or executes no actions at all

Mechanisms for Ensuring Atomicity

- Logging
 - DBMS logs all actions so that it can undo the actions of aborted transactions
- Shadowing
 - DBMS makes copies of pages and transactions make changes to those copies
 - Only when the transaction commits the page is made visible to others

We take \$100 out of Andy's account but then there is a power failure before we transfer it to his bookie.

When the database comes back on-line, what should be the correct state of Andy's account?

Consistency

Database Consistency

- Data in the DBMS is accurate in modeling the real world and follows integrity constraints

Transaction Consistency

- If the database is consistent before the transactions starts (running alone), it also will be after
- Transaction consistency is the application's responsibility
 - We won't discuss this further...

Users submit transactions, and each transaction executes as if it was running by itself

Concurrency is achieved by DBMS, which interleaves actions (reads/writes of DB objects) of various transactions

How do we achieve this? → Many methods - two main categories

- **Pessimistic** – Don't let problems arise in the first place
- **Optimistic** – Assume conflicts are rare, deal with them after they happen

All results of successful transactions have to be made persistent

Example

- If a flight booking reports that a seat has successfully been booked, then the seat will remain booked even if the system crashes

Durability can be achieved by flushing the transaction's log records to non-volatile storage before acknowledging commitment

- Write-Ahead Log (WAL): Record the changes made to the database in a log before the change is made



Isolation

Example

Two transactions

- T_1 transfers \$100 from B's account to A's
- T_2 credits both accounts with 6% interest

Assume at first A and B each have \$1000

What are the legal outcomes of running T_1 and T_2 ?

- Many! But $A+B$ should be: $\$2000 \times 1.06 = \2120
- There is no guarantee that T_1 will execute before T_2 or vice-versa, if both are submitted together
- But, the net effect must be equivalent to these two transactions running serially in some order
- Legal outcomes
 - $A=1166, B=954$
 - $A=1160, B=960$
- The outcome depends on whether T_1 executes before T_2 or vice versa

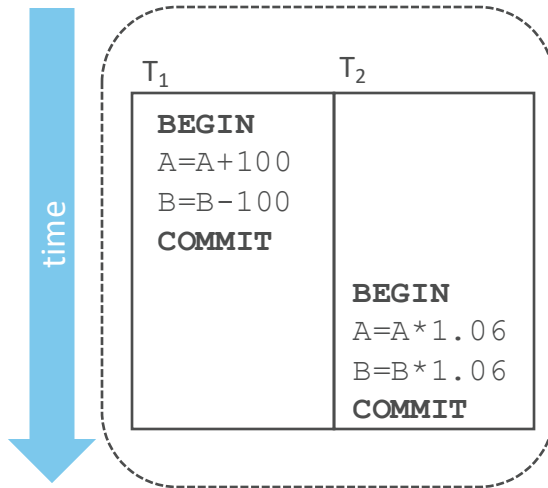
T_1	T_2
BEGIN	BEGIN
$A=A+100$	$A=A \times 1.06$
$B=B-100$	$B=B \times 1.06$
COMMIT	COMMIT

Example


Serial Execution

- T_1 transfers \$100 from B's account to A's
- T_2 credits both accounts with 6% interest

Schedule



$A=1166, B=954$



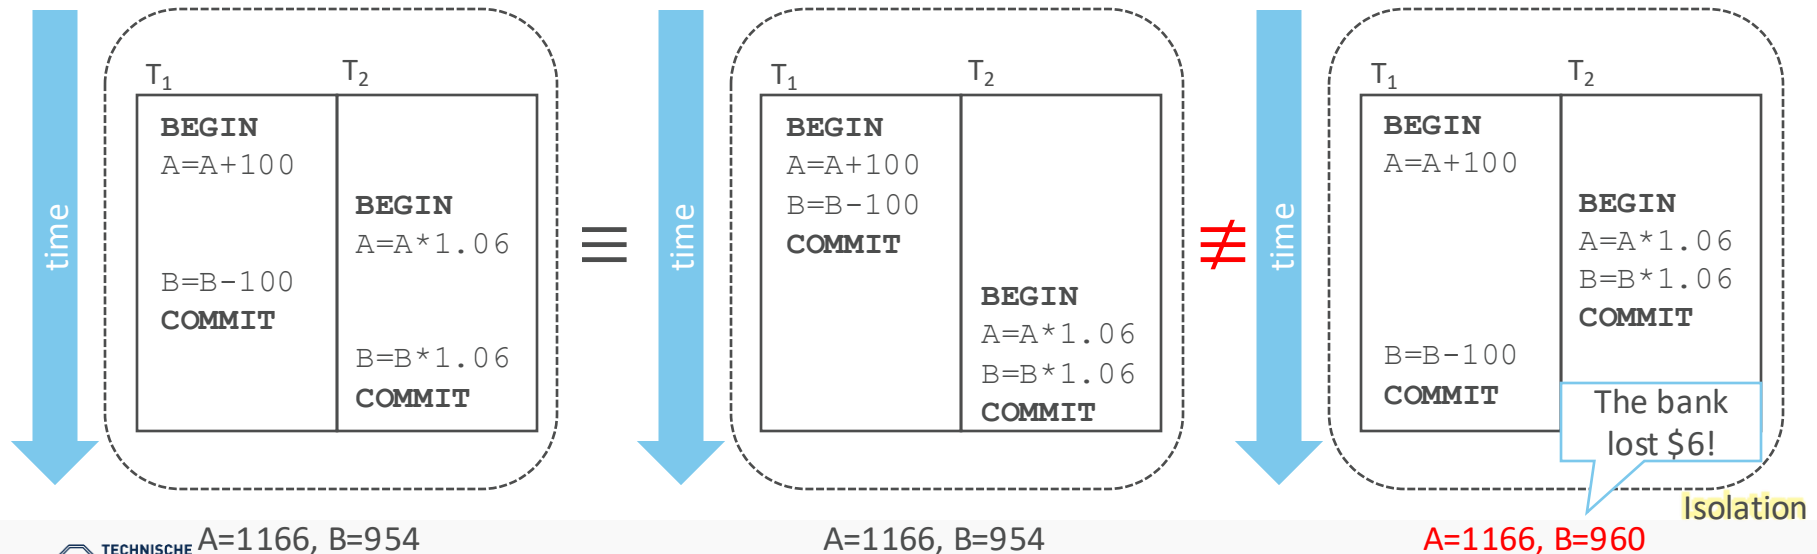
$A=1160, B=960$

Interleaving Transactions

We can also interleave the transactions in order to maximize concurrency

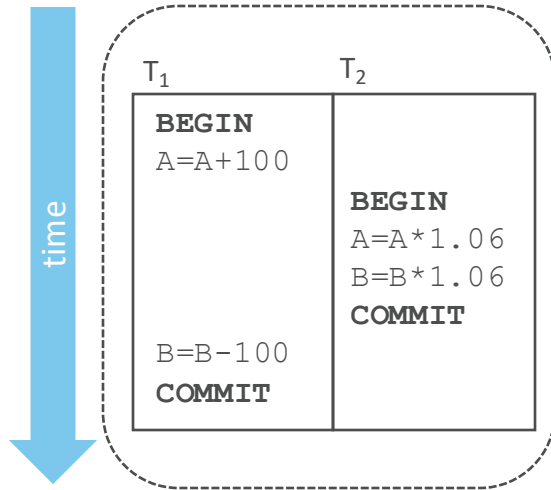
- Slow disk/network I/O
- Multi-core CPUs

Interleaving Examples




Schedule versus DBMS View

Schedule



DBMS's View



How do we judge that a schedule is correct?

- If it is equivalent to some serial execution

Serial Schedule

- A schedule that does not interleave the actions of different transactions

Equivalent Schedules

- For any database state, the effect of executing the first schedule is identical to the effect of executing the second schedule → **no matter what the arithmetic operations are!**

Serializable Schedule

- A schedule that is equivalent to **some serial execution** of the transactions
- Note: If each transaction preserves consistency, every serializable schedule preserves consistency

Serializability

- Less intuitive notion of correctness compared to transaction initiation time or commit order, but it provides the DBMS with significant additional **flexibility** in scheduling operations.

Isolation




Anomalies

Interleaved Execution Anomalies


Write-Read conflicts (W-R)

- Reading Uncommitted Data → “Dirty Reads”



Read-Write Conflicts (R-W)


- Unrepeatable Reads



Interleaved Execution Anomalies

Write-Write Conflicts (W-W)

- Overwriting uncommitted data → lost update



How could you guarantee that all resulting schedules are correct (i.e., serializable)?

→ Use locks!



Locking

Locks

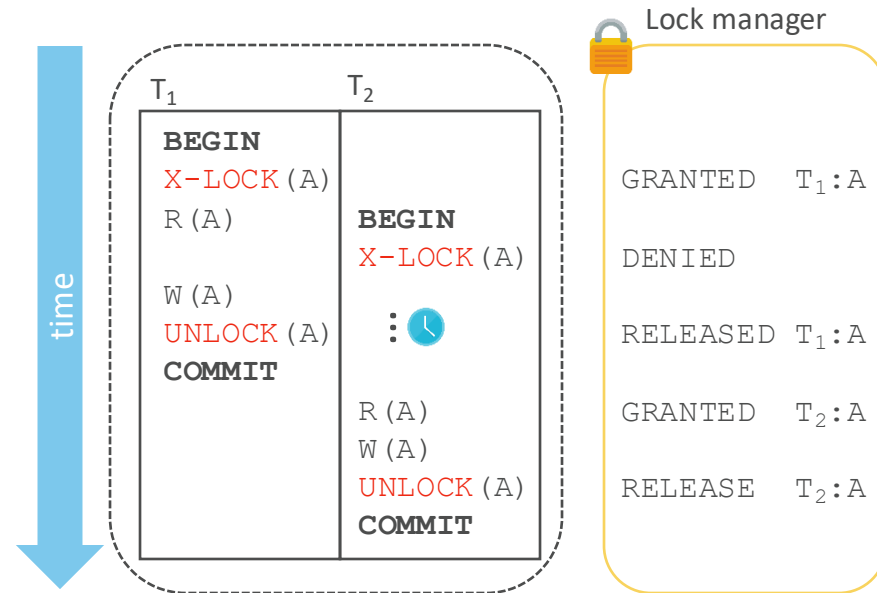
Basic Lock Types

- S-LOCK – Shared Locks (reads)
- X-LOCK – Exclusive Locks (writes)

Compatibility Matrix

	shared	exclusive
shared	✓	✗
exclusive	✗	✗

Example



Locks (2)

Executing with Locks

- Transactions request locks (or upgrades)
- Lock manager grants or blocks requests
- Transactions release locks
- Lock manager updates lock-table

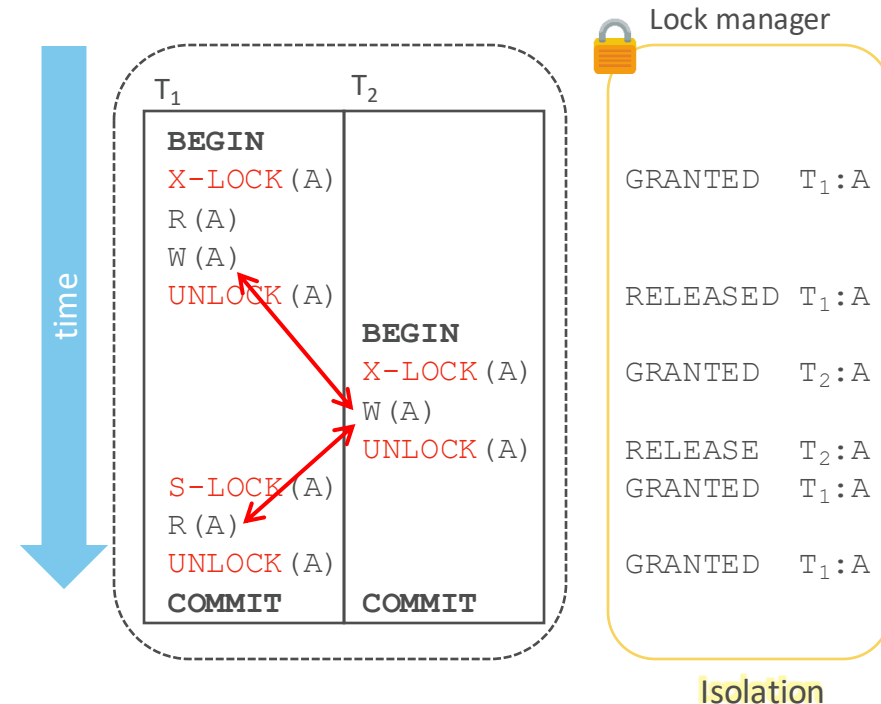
→ But this is not enough...

- We need to use a well-defined protocol that ensures that transactions execute correctly

One possible solution

- Two-Phase Locking (2PL)

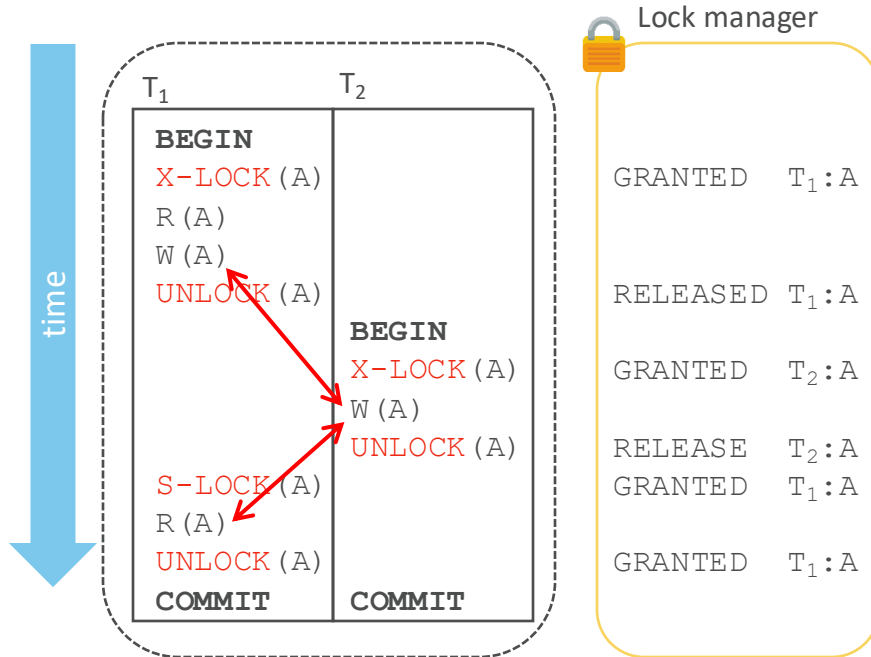
Example



Two-Phase locking

Motivation

- Example




Phase 1: Growing

- Each transaction requests the locks that it needs from the DBMS's lock manager
- The lock manager grants/denies lock requests

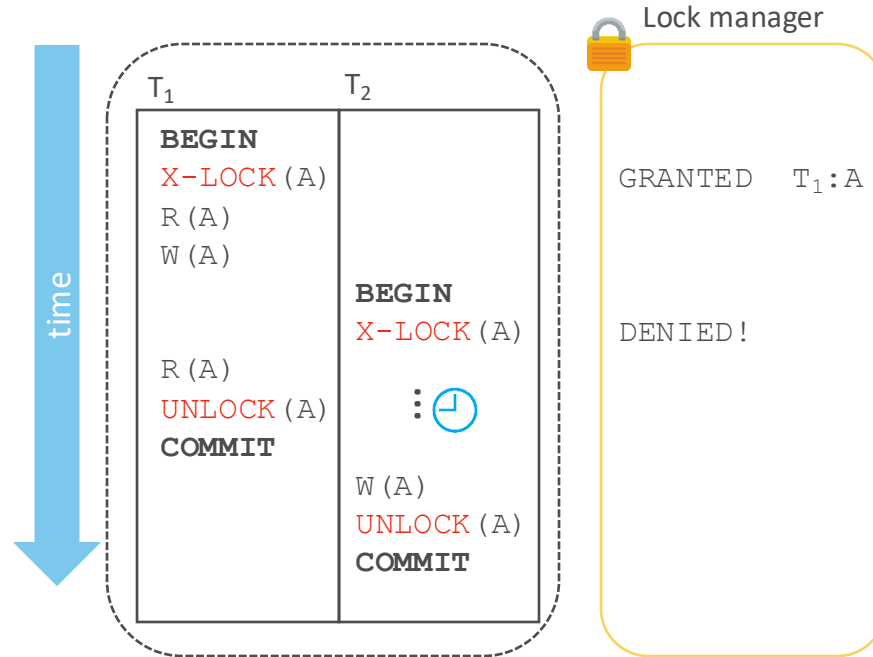
Phase 2: Shrinking

- The transaction is allowed to only release locks that it previously acquired
- It cannot acquire new locks

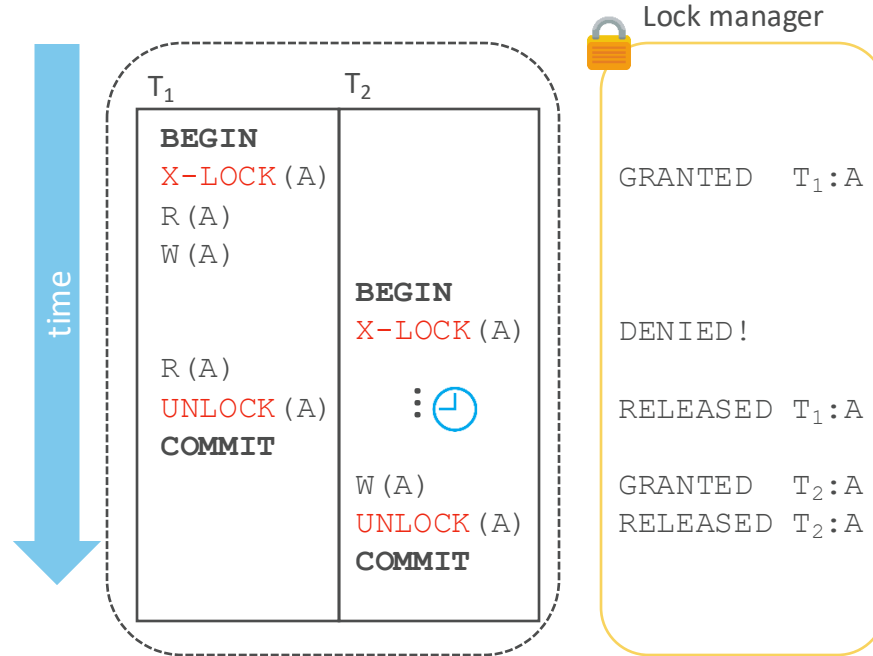
Two-Phase Locking



Example



Example



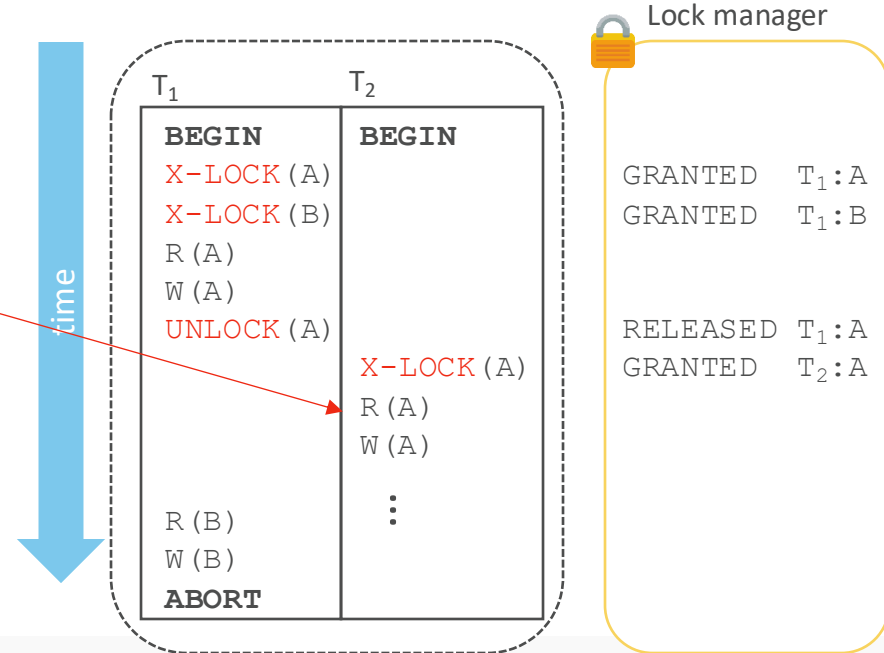
Two-Phase Locking

2PL on its own is...

- ...sufficient to guarantee conflict serializability (i.e., schedules whose precedence graph is acyclic)
- but, it is subject to **cascading aborts**

Example

- This is a permissible schedule in 2PL, but we have to abort T_2 too
- This is all wasted work!



Strict Two-Phase Locking

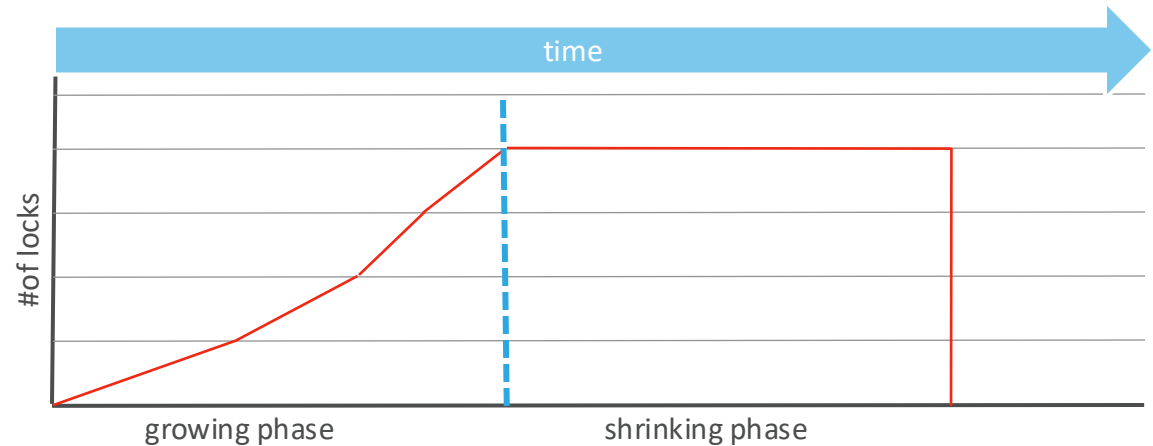
A transaction releases its write locks only after it has ended (committed or aborted)

A schedule is strict if...

- ...a value written by a transaction is not read or overwritten by other transaction until that transaction finishes

Advantages

- Does **not** incur **cascading aborts**
- Aborted transactions can be undone by just restoring original values of modified tuples





Deadlocks

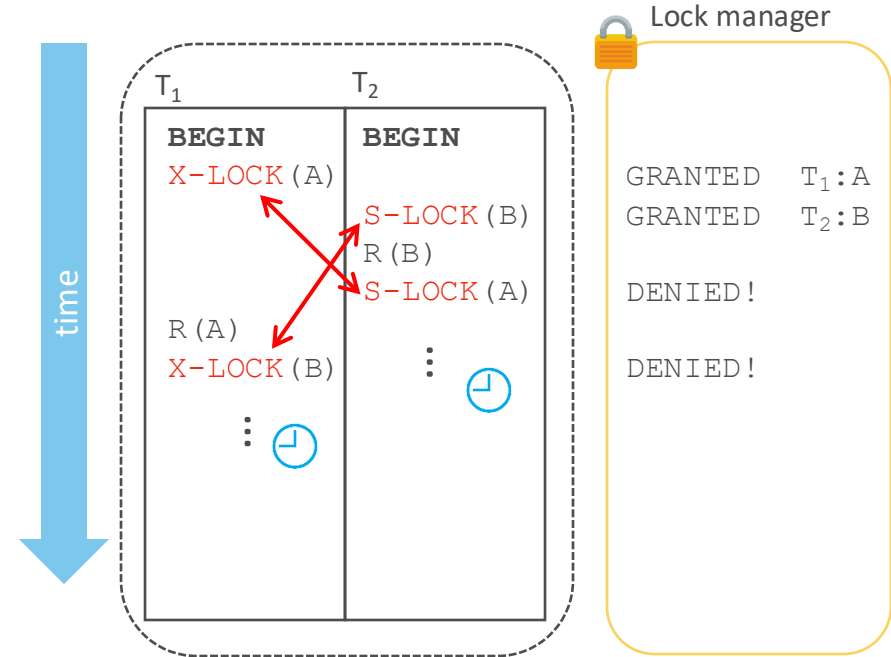
Deadlock

Deadlock

- Cycle of transactions waiting for locks to be released by each other

Two ways of dealing with deadlocks

- Deadlock **detection**
- Deadlock **prevention**



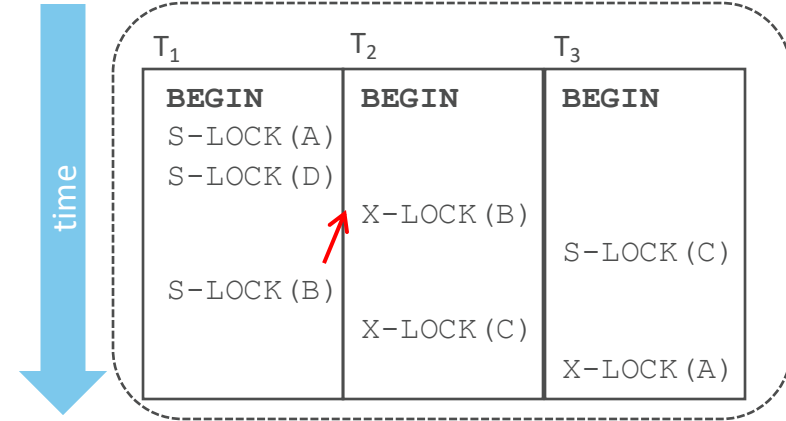
Deadlock Detection

The DBMS creates a waits-for graph


- Nodes are transactions
- Edge from T_i to T_j if T_i is waiting for T_j to release a lock

The system periodically check for cycles in waits-for graph

Schedule



Waits-for graph



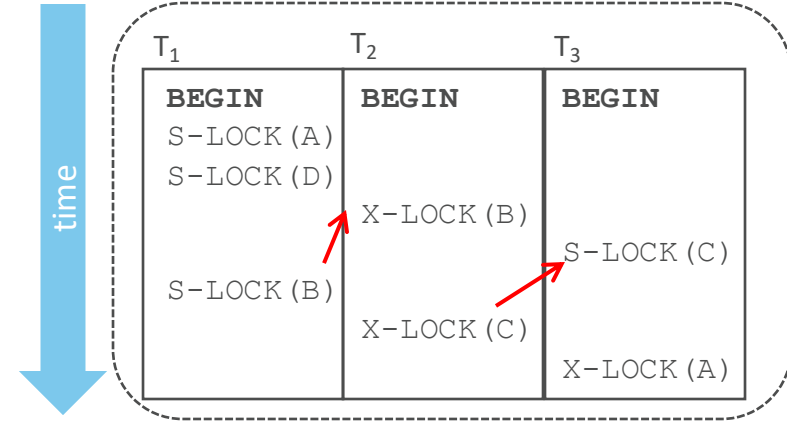
Deadlock Detection

The DBMS creates a waits-for graph


- Nodes are transactions
- Edge from T_i to T_j if T_i is waiting for T_j to release a lock

The system periodically check for cycles in waits-for graph

Schedule



Waits-for graph




Deadlock Detection

The DBMS creates a waits-for graph


- Nodes are transactions
- Edge from T_i to T_j if T_i is waiting for T_j to release a lock

The system periodically check for cycles in waits-for graph

Schedule



Waits-for graph




Deadlock Handling

What do we do if a deadlock occurs?

- Select a “victim” and rollback it back to break the deadlock

Which one do we choose?

- Decide by considering
 - age (lowest timestamp)
 - progress (least/most operations issued)
 - # of items already locked
 - # of transaction that we have to rollback with it
- We also should consider the # of times a transaction has been restarted in the past



When a transaction tries to acquire a lock that is held by another transaction, kill one of them to prevent a deadlock


- No waits-for graph or detection algorithm

Assign priorities based on timestamps

- Older \rightarrow higher priority (e.g., $T_1 > T_2$)
- Two different prevention policies
 - Wait-Die: If T_1 has higher priority, T_1 waits for T_2 ; otherwise T_1 aborts (“old wait for young”)
 - Wound-Wait: If T_1 has higher priority, T_2 aborts; otherwise T_1 waits (“young wait for old”)

Deadlock Prevention

Examples




Wait-Die

- T₁ waits

Wound-Wait

- T₂ aborted



Wait-Die

- T₂ aborted

Wound-Wait

- T₂ waits

Summary

Transaction

- Sequence of operations
- Serial vs. interleaving execution

ACID

- Atomicity
- Consistency
- Isolation
- Durability

Isolation

- Anomalies
- Locking
- Deadlocks